

AI Response Evaluation Rubric

Project: AI Response Evaluation Study

Evaluator: ANJALI

1. Objective

The objective of this rubric is to evaluate AI-generated responses consistently across multiple quality dimensions. The evaluation focuses on accuracy, relevance, completeness, clarity, instruction following, and overall response quality.

The evaluation focuses on:

1. Accuracy
2. Relevance
3. Completeness
4. Clarity
5. Instruction Following
6. Overall Quality

Each criterion is scored from **1 to 5**. Evaluators should provide evidence-based comments when assigning scores.

2. Evaluation Scale

Score 1 — Very Poor

The response has serious problems and fails to satisfy the user's request.

Score 2 — Poor

The response has significant weaknesses and requires substantial improvement.

Score 3 — Acceptable

The response addresses the main request but has noticeable weaknesses or omissions.

Score 4 — Good

The response is strong overall with only minor issues.

Score 5 — Excellent

The response fully satisfies the request and has no significant identifiable problems.

3. Evaluation Criteria

A. Accuracy

Evaluation Question

Are the factual claims in the response correct?

Score	Definition
1	Contains major factual errors or misinformation that substantially affect the answer.
2	Contains multiple factual errors or significant inaccuracies.
3	Mostly accurate but contains minor factual errors, questionable claims, or inaccuracies.
4	Accurate overall with only minor issues or omissions.
5	Fully accurate with no identifiable factual errors.

Evaluator should check:

- Factual correctness
- Numerical correctness
- Logical factual claims
- Dates and historical information
- Definitions
- Scientific information
- Unsupported claims

Example

Prompt: What is the capital of Australia?

Response: Sydney is the capital of Australia.

Score: 1/5

Reason: The response contains a major factual error. Canberra is the capital of Australia.

B. Relevance

Evaluation Question

Does the response directly address the user's request without unnecessary information?

Score	Definition
1	Does not address the user's question or task.
2	Addresses the topic only minimally and contains substantial irrelevant information.
3	Partially addresses the request but includes noticeable irrelevant information.
4	Directly addresses the request with minor unnecessary information.
5	Directly addresses the request and remains focused throughout.

Evaluator should check:

- Does it answer the actual question?
 - Does it stay on topic?
 - Does it avoid unnecessary tangents?
 - Does it provide information relevant to the user's goal?
-

C. Completeness

Evaluation Question

Does the response adequately address all important parts of the user's request?

Score	Definition
1	Almost completely fails to address the requested information.
2	Addresses only a small portion of the request.
3	Addresses the main request but misses important information.
4	Covers nearly all important aspects with minor omissions.
5	Fully addresses all important aspects of the request.

Example

Prompt:

"Give me three advantages of exercise and explain each one."

If the AI gives three advantages but does not explain them:

Completeness: 3/5

Reason:

The response provides the requested three advantages but does not adequately explain each one.

D. Clarity

Evaluation Question

Is the response easy to understand, logically organized, and appropriately written?

Score	Definition
1	Extremely difficult to understand, incoherent, or severely poorly organized.
2	Frequently confusing, poorly structured, or difficult to follow.

- | | |
|---|--|
| 3 | Generally understandable but contains noticeable clarity or organization problems. |
| 4 | Clear and well organized with minor issues. |
| 5 | Very clear, logical, well organized, and easy to understand. |

Evaluator should consider:

- Grammar
 - Sentence structure
 - Organization
 - Logical flow
 - Vocabulary
 - Readability
 - Appropriate level of explanation
-

E. Instruction Following

Evaluation Question

Did the AI follow all explicit instructions in the user's prompt?

Score	Definition
1	Ignores most or all explicit instructions.
2	Follows only a small portion of the instructions.
3	Follows some instructions but misses one or more important requirements.
4	Follows nearly all instructions with minor deviations.
5	Follows all explicit instructions correctly.

Important

Check requirements individually.

For example, if the prompt says:

"Write exactly 100 words, use five bullet points, and provide one example."

Check:

- Word count
- Five bullet points
- Exactly one example
- Subject matter
- Any required formatting

Do not give a high score simply because the response is well written.

A response can be excellent in quality but still receive a low **Instruction Following** score.

F. Overall Quality

Evaluation Question

How useful is the response to the user overall?

Score	Definition
1	Unusable; major problems prevent the response from fulfilling the request.
2	Poor; significant problems require substantial correction.
3	Acceptable; useful but requires noticeable improvement.
4	Good; useful and mostly meets the user's needs.
5	Excellent; accurate, relevant, complete, clear, and highly useful.

Important

Overall Quality should reflect your **overall judgment**, but it should not simply be the mathematical average of the other five scores.

For example:

A response could receive:

- Accuracy: 5
- Relevance: 5
- Completeness: 3
- Clarity: 5
- Instruction Following: 2

Its Overall Quality might reasonably be **3 or 4**, depending on how important the missed instructions were.

Explain your decision in the comments.

4. Evaluation Guidelines

At the end of your rubric, create:

Evaluation Guidelines

Add these rules:

1. Evaluate the response, not the AI model.

Scores should be based on the actual response being evaluated.

2. Use evidence.

When assigning a low or high score, provide a brief explanation.

3. Evaluate each criterion independently.

A response can be clear but factually incorrect, or accurate but incomplete.

4. Do not give higher scores because the answer sounds confident.

Verify factual claims when necessary.

5. Apply the same standards to every response.

6. Do not change the rubric while evaluating the dataset.

If a change is necessary, document the change and apply the revised rule consistently.

That last rule is particularly important.

5. Factual Verification

For factual questions, the evaluator should verify important factual claims using reliable sources when the evaluator is uncertain about the accuracy of the response.

Examples of useful sources include:

- Government websites
- Universities
- Scientific institutions
- Official organizations
- Academic publications
- Reputable reference sources

This prevents you from simply assuming an AI answer is correct.

6. Evaluator Comment Guidelines

Comments should explain **why** a score was given.

Weak comment

Good answer.

Better comment

The response directly answers the question and provides an accurate explanation, but it omits an important example requested by the user.

Strong comment

The response correctly explains the main concept and remains focused on the question. However, it does not address the second requirement in the prompt and therefore receives a lower instruction-following score.

Comments should be:

- Specific
- Concise
- Evidence-based
- Related to the evaluation criterion

7. Evaluation Independence

The evaluator should evaluate the **actual response**, not the reputation of the AI model.

Do not assume that one model is automatically better than another.

For every response:

1. Read the prompt.
 2. Read the response.
 3. Apply the rubric.
 4. Verify factual claims when necessary.
 5. Assign scores.
 6. Write a brief justification.
-

8. Consistency Rules

Apply the same standards to every response.

Rule 1 — Use the same rubric

Do not change the scoring criteria for different AI models.

Rule 2 — Evaluate independently

Do not give a response a higher score because you personally prefer the model that generated it.

Rule 3 — Use evidence

Base scores on identifiable characteristics of the response.

Rule 4 — Separate criteria

Evaluate each criterion independently.

For example:

A response can be:

- Clear but inaccurate
- Accurate but incomplete

- Relevant but poorly written
- Complete but fail specific instructions

Rule 5 — Document changes

If the rubric must be changed during the project, document the change and apply the revised rule consistently where appropriate.

9. Special Guidelines for Your 10 Prompt Categories

General Knowledge

Focus on:

- Factual accuracy
- Relevance
- Completeness
- Clear explanation

Science

Focus on:

- Scientific accuracy
- Correct terminology
- Logical explanation
- Avoidance of misleading claims

History

Focus on:

- Historical accuracy
- Dates and events where relevant
- Cause-and-effect relationships
- Avoidance of unsupported claims

Technology

Focus on:

- Technical accuracy
- Appropriate explanations
- Correct terminology
- Avoidance of outdated or misleading information

Business

Focus on:

- Correct business concepts
- Logical reasoning
- Practical relevance
- Clear explanations

Education

Focus on:

- Accuracy
- Clarity
- Practical usefulness
- Appropriate explanations

Writing

Focus on:

- Grammar
- Structure
- Tone
- Clarity
- Relevance
- Requested length

Reasoning

Focus on:

- Correct conclusion
- Logical reasoning
- Mathematical accuracy where applicable
- Whether the explanation supports the conclusion

Instruction Following

Focus heavily on:

- Every explicit requirement
- Word count
- Number of items
- Required format
- Required structure
- Required content

Ambiguous Questions

Focus on whether the AI:

- Recognizes missing information
 - Avoids unjustified assumptions
 - Asks useful clarifying questions when appropriate
 - Explains relevant factors
 - Provides a conditional or contextual answer when appropriate
-

10. Recommended Evaluation Record

For every response, record:

Field	Value
Prompt ID	Example: GK-01
Category	General Knowledge
AI Model	Model name
Accuracy	1–5
Relevance	1–5
Completeness	1–5
Clarity	1–5

Instruction Following 1–5

Overall Quality 1–5

Evaluator Comments Short explanation

Verification Source If required

11. Quality-Control Checklist

Before finalizing an evaluation, ask:

- Did I read the complete prompt?
 - Did I read the complete AI response?
 - Did I evaluate accuracy separately from clarity?
 - Did I check whether all instructions were followed?
 - Did I check important factual claims?
 - Did I avoid judging the AI model itself?
 - Did I provide evidence for unusual or important scores?
 - Did I apply the same standards used for other responses?
 - Did I record my reasoning clearly?
 - Did I avoid changing the rubric just because a particular response was difficult to score?
-

12. Final Scoring Framework

Each response receives a score for six dimensions:

Accuracy: 1–5

Relevance: 1–5

Completeness: 1–5

Clarity: 1–5

Instruction Following: 1–5

Overall Quality: 1–5

The first five dimensions provide structured evidence about response quality, while Overall Quality represents the evaluator's final judgment.

Maximum score per response: 30 points.

However, the project should primarily analyze the **individual dimensions**, rather than relying only on a combined total score.
